

Practice Problems in R

Peter Sun

September 30-31, 2021

Contents

1	Practice 1: Generalized Boosted Regression and Propensity Score Weighting	2
1.1	Problem 1: Generalized Boosted Regression	2
1.1.1	Load Packages	2
1.1.2	Load and Randomly Shuffle Data	2
1.1.3	Generate Propensity Scores	2
1.1.4	Histograms of Propensity Scores	3
1.1.5	Boxplots of Propensity Scores	4
1.2	Problem 2: Propensity Score Weighting	5
1.2.1	Estimate ATE and ATT Weights	5
1.2.2	Outcome Analysis with ATE and ATT Weights	5
1.2.3	Check Imbalance	7
1.2.4	Alternative Solution with the WeightIt Package	9
2	Practice 2: Matching Estimators	12
2.1	Load Data	12
2.2	Breusch-Pagan Test for Heteroskedasticity	12
2.3	Matching Estimators	13
2.3.1	Define Outcome (Y), Treatment Index (Tr), and Variables to Match On (X)	13
2.3.2	Define Function for Matching	13
2.3.3	Get All Estimators	13
3	Appendix: Replicating Stata's Robust Standard Errors	14

1 Practice 1: Generalized Boosted Regression and Propensity Score Weighting

1.1 Problem 1: Generalized Boosted Regression

1.1.1 Load Packages

```
library(tidyverse)
library(haven)
library(sjlabelled)
library(lmtest)
library(gbm)
library(modelr)
library(broom)
library(sandwich)
library(cobalt)
library(WeightIt)
library(Matching)
library(kableExtra)
select <- dplyr::select
```

1.1.2 Load and Randomly Shuffle Data

```
d <- read_dta("data/ldw_exper.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble()
set.seed(1000)
d2 <- d %>%
  add_column(runif = runif(nrow(.))) %>%
  arrange(runif)
```

1.1.3 Generate Propensity Scores


```
set.seed(1000)
m1 <- gbm::gbm(formula = t ~ age + educ + black + hisp + married + re74 +
                re75 + u74 + u75,
               data = d2,
               distribution = "bernoulli",
               n.trees = 1000,
               train.fraction = 0.8,
               interaction.depth = 4,
               shrinkage = 0.0005)

# Estimate Propensity Scores and Obtain Summary Statistics
d3 <- d2 %>%
  modelr::add_predictions(m1, var = "psb", type = "response")
summary(d3$psb)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.3394  0.3855  0.4006  0.4108  0.4346  0.5214
```

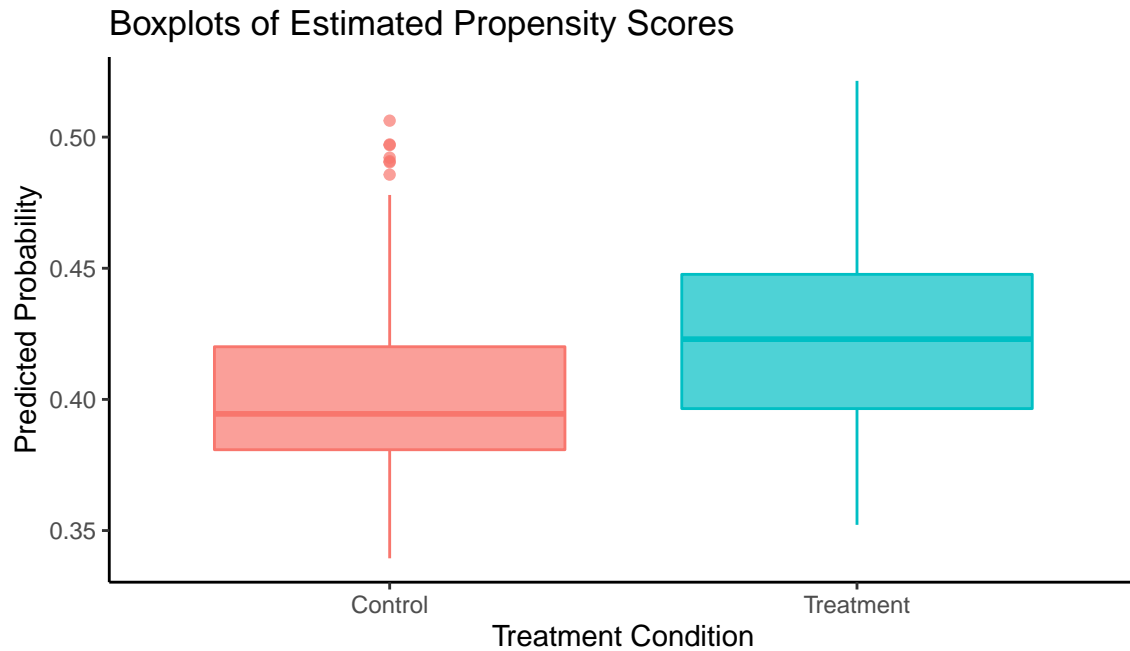
1.1.4 Histograms of Propensity Scores

```
d3 %>%  
  mutate(t = factor(t, labels = c("Control", "Treatment"))) %>%  
  ggplot(aes(x = psb, color = t)) +  
  theme_classic() +  
  geom_histogram(aes(fill = t), alpha = 0.6, bins = 13) +  
  geom_density(size = 1) +  
  labs(x = "Predicted Probability", y = "Density",  
       title = "Histograms of Estimated Propensity Scores") +  
  theme(legend.position = "none") +  
  facet_wrap(~ t)
```



1.1.5 Boxplots of Propensity Scores

```
d3 %>%  
  mutate(t = factor(t, labels = c("Control", "Treatment"))) %>%  
  ggplot(aes(x = t, y = psb, color = t, fill = t)) +  
  theme_classic() +  
  geom_boxplot(alpha = 0.7) +  
  labs(x = "Treatment Condition",  
       y = "Predicted Probability",  
       title = "Boxplots of Estimated Propensity Scores") +  
  theme(legend.position = "none")
```



1.2 Problem 2: Propensity Score Weighting

1.2.1 Estimate ATE and ATT Weights

```
d4 <- d3 %>%
  mutate(ate_w = ifelse(t == 0, 1/(1-psb), 1/psb),
         att_w = ifelse(t == 0, psb/(1-psb), 1))

# Import Stata-generated weights to replicate Stata results
stata_weights <- read_dta("data/ldw1.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble() %>%
  mutate(stata_ate_w = ifelse(t == 0, 1/(1-psb), 1/psb),
         stata_att_w = ifelse(t == 0, psb/(1-psb), 1)) %>%
  select(id, stata_ate_w, stata_att_w)
d5 <- d4 %>%
  arrange(id) %>%
  left_join(stata_weights, by = "id")
```

1.2.2 Outcome Analysis with ATE and ATT Weights

```
# Define outcome formula
f = as.formula(re78 ~ t + age + educ + black + hisp + married + re74 + re75 +
              u74 + u75)

# Weighted OLS with R-Generated Propensity Scores
m2 <- lm(f, data = d5, weights = ate_w)
tidy(lmtest::coeftest(m2, vcov. = vcovHC(m2, "HC1"))) %>% filter(term == "t") # ATE

## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>      <dbl>      <dbl>      <dbl>  <dbl>
## 1 t          1.65        0.655        2.52  0.0122

m3 <- lm(f, data = d5, weights = att_w)
tidy(lmtest::coeftest(m3, vcov. = vcovHC(m3, "HC1"))) %>% filter(term == "t") # ATT

## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>      <dbl>      <dbl>      <dbl>  <dbl>
## 1 t          1.72        0.663        2.60  0.00965

# Weighted OLS with Stata-Generated Propensity Scores (Identical Results)
m2.stata <- lm(f, data = d5, weights = stata_ate_w)
tidy(lmtest::coeftest(m2.stata, vcov. = vcovHC(m2.stata, "HC1"))) %>%
  filter(term == "t") # ATE

## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>      <dbl>      <dbl>      <dbl>  <dbl>
## 1 t          1.63        0.656        2.48  0.0135

m3.stata <- lm(f, data = d5, weights = stata_att_w)
tidy(lmtest::coeftest(m3.stata, vcov. = vcovHC(m3.stata, "HC1"))) %>%
  filter(term == "t") # ATT
```

```
## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>      <dbl>      <dbl>      <dbl>  <dbl>
## 1 t          1.70        0.665        2.55  0.0111
```

1.2.3 Check Imbalance

See the Appendix for the custom function `robustse()` that is used to replicate the robust standard errors in Stata.

```
# Function to Check Imbalance
check_bal <- function(var, weight, type) {
  if(type == "categorical") {
    m <- glm(as.formula(paste0(var, "~t")),
             family = quasibinomial,
             data = d5,
             weights = weight
            )
    m %>%
      tidy() %>%
      mutate(odds.ratio = exp(estimate), variable = var) %>%
      mutate(or.se = robustse(m, coef = "odds.ratio")[,2]) %>%
      mutate(statistic = robustse(m, coef = "odds.ratio")[,3]) %>%
      mutate(p.value = robustse(m, coef = "odds.ratio")[,4]) %>%
      select(variable, term, odds.ratio, or.se, statistic, p.value)
  } else if(type == "continuous") {
    m <- lm(as.formula(paste0(var, "~t")),
            data = d5,
            weights = weight)
    lmtest::coeftest(m, vcov. = vcovHC(m, "HC1")) %>%
      tidy() %>%
      add_column(var, .before = "term")
  }
}

format_bal <- function(df) {
  df %>%
    filter(term != "(Intercept)") %>%
    kbl(booktabs = T, digits = 7) %>%
    kable_styling(position = "center") %>%
    kable_styling(latex_options = c("striped", "HOLD_position"))
}
```

Categorical Variables

```
cat_vars <- c("black", "hisp", "married", "u74", "u75")
format_bal(map_dfr(cat_vars, check_bal, d5$stata_ate_w, "categorical"))
```

variable	term	odds.ratio	or.se	statistic	p.value
black	t	1.0997119	0.2872935	0.3638291	0.7159856
hisp	t	0.5680117	0.2106900	-1.5248703	0.1272914
married	t	1.2388627	0.3163250	0.8388736	0.4015402
u74	t	0.8907275	0.1929677	-0.5341418	0.5932434
u75	t	0.7896914	0.1590241	-1.1725044	0.2409946

```
format_bal(map_dfr(cat_vars, check_bal, d5$stata_att_w, "categorical"))
```

variable	term	odds.ratio	or.se	statistic	p.value
black	t	1.0785012	0.2815619	0.2894740	0.7722186
hisp	t	0.5688118	0.2106633	-1.5234109	0.1276559
married	t	1.2497739	0.3193590	0.8725380	0.3829149
u74	t	0.8616888	0.1868079	-0.6866516	0.4923023
u75	t	0.7639094	0.1539378	-1.3364198	0.1814121

Continuous Variables

```
cont_vars <- c("age", "educ", "re74", "re75")
format_bal(map_dfr(cont_vars, check_bal, d5$stata_ate_w, "continuous"))
```

var	term	estimate	std.error	statistic	p.value
age	t	0.5085618	0.6801859	0.7476806	0.4550495
educ	t	0.1629619	0.1762747	0.9244774	0.3557411
re74	t	-0.1742469	0.4893238	-0.3560973	0.7219372
re75	t	0.1288075	0.2977874	0.4325485	0.6655533

```
format_bal(map_dfr(cont_vars, check_bal, d5$stata_att_w, "continuous"))
```

var	term	estimate	std.error	statistic	p.value
age	t	0.5532783	0.6945403	0.7966107	0.4261038
educ	t	0.1874483	0.1813631	1.0335526	0.3019093
re74	t	-0.0963895	0.5091306	-0.1893218	0.8499273
re75	t	0.1832357	0.3088324	0.5933178	0.5532713

Similar results can be obtained using the R-generated propensity score weights:

With R-generated weights

```
format_bal(map_dfr(cat_vars, check_bal, d5$ate_w, "categorical"))
format_bal(map_dfr(cat_vars, check_bal, d5$att_w, "categorical"))
format_bal(map_dfr(cont_vars, check_bal, d5$ate_w, "continuous"))
format_bal(map_dfr(cont_vars, check_bal, d5$att_w, "continuous"))
```


1.2.4 Alternative Solution with the WeightIt Package


Use GBM to estimate ATE and ATT:

```
set.seed(1000)
w1.out <- WeightIt::weightit(
  formula = t ~ age + educ + black + hisp + married + re74 +
    re75 + u74 + u75,
  data = d2,
  method = "gbm",
  distribution = "bernoulli",
  stop.method = "es.mean",
  n.trees = 1000,
  nTrain = 0.8 * nrow(d2),
  interaction.depth = 4,
  shrinkage = 0.0005,
  estimand = "ATE")


set.seed(1000)
w2.out <- WeightIt::weightit(
  formula = t ~ age + educ + black + hisp + married + re74 +
    re75 + u74 + u75,
  data = d2,
  method = "gbm",
  distribution = "bernoulli",
  stop.method = "es.mean",
  n.trees = 1000,
  nTrain = 0.8 * nrow(d2),
  interaction.depth = 4,
  shrinkage = 0.0005,
  estimand = "ATT")
```

Assess balance with the cobalt package:

```
cobalt::love.plot(w1.out, thresholds = c(m = .1), binary = "std", abs = T) +
  labs(title = "Covariate Balance (ATE)")
```



```
cobalt::love.plot(w2.out, thresholds = c(m = .1), binary = "std", abs = T) +
  labs(title = "Covariate Balance (ATT)")
```



For the outcome analysis, the ATE and ATT weights can be obtained with `w1.out$weights` (ATE) and `w2.out$weights` (ATT):

```
m2.weightit <- lm(f, data = d2, weights = w1.out$weights)
tidy(lmtest::coefest(m2.weightit, vcov. = vcovHC(m2.weightit, "HC1"))) %>%
  filter(term == "t")
```

```
## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>    <dbl>    <dbl>    <dbl>  <dbl>
## 1 t        1.57      0.646      2.43  0.0155
```

```
m3.weightit <- lm(f, data = d2, weights = w2.out$weightit)
tidy(lmtest::coefest(m3.weightit, vcov. = vcovHC(m3.weightit, "HC1"))) %>%
  filter(term == "t")
```

```
## # A tibble: 1 x 5
##   term estimate std.error statistic p.value
##   <chr>     <dbl>     <dbl>     <dbl>   <dbl>
## 1 t         1.67      0.662      2.53  0.0119
```

2 Practice 2: Matching Estimators

2.1 Load Data

```
p2.d <- read_dta("data/prac2.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble()
```

2.2 Breusch-Pagan Test for Heteroskedasticity

The homoscedasticity assumption is not valid (e.g., p-value of the test for `age97` is $< .05$), indicating that the conditional variance of the outcome variable was not constant across levels of child's age, therefore a robust estimation of variance is warranted.

```
p2.m0 <- lm(lwss97 ~ kuse + male + black + age97 + pcged97 + mratio96 + pcg_adc,
            data = p2.d)
get_bpctest <- function(data, lm.model, var) {
  b <- lmtest::bptest(lm.model, as.formula(paste0("~", var)),
                      data = data, studentize = F)
  return(tibble(variable = var, statistic = b$statistic,
                df = b$parameter, p.value = b$p.value))
}
map_dfr(c("kuse", "male", "black", "age97", "pcged97", "mratio96", "pcg_adc"),
        get_bpctest, data = p2.d, lm.model = p2.m0) %>%
  kbl(booktabs = T, digits = 2, linesep = "",
      caption = "Results of Breusch-Pagan Tests for Heteroskedasticity") %>%
  kable_styling(position = "center") %>%
  kable_styling(latex_options = c("striped", "HOLD_position"))
```

Table 1: Results of Breusch-Pagan Tests for Heteroskedasticity

variable	statistic	df	p.value
kuse	1.78	1	0.18
male	0.86	1	0.35
black	1.15	1	0.28
age97	8.55	1	0.00
pcged97	4.43	1	0.04
mratio96	6.85	1	0.01
pcg_adc	0.60	1	0.44

2.3 Matching Estimators

2.3.1 Define Outcome (Y), Treatment Index (Tr), and Variables to Match On (X)

```
Y <- p2.d$lwss97
Tr <- p2.d$kuse
X <- select(p2.d, male, black, age97, pcged97, mratio96, pcg_adc)
```

2.3.2 Define Function for Matching

```
get_match <- function(estimand, sample) {
  m <- Matching::Match(Y = Y, Tr = Tr, X = X, M = 4, BiasAdjust = T, Var.calc = 4,
    estimand = estimand, sample = sample)
  return(list(
    est = m$est[,1],
    se = m$se,
    t.stat = m$est[,1]/m$se,
    p = (1 - pnorm(abs(m$est[,1]/m$se))) * 2
  ))
}
```

2.3.3 Get All Estimators

```
tribble(
  ~estimator, ~estimand, ~sample,
  "SATE", "ATE", T,
  "PATE", "ATE", F,
  "SATT", "ATT", T,
  "PATT", "ATT", F,
  "SATC", "ATC", T,
  "PATC", "ATC", F
) %>%
rowwise() %>%
mutate(match = list(get_match(estimand, sample))) %>%
tidyr::unnest_wider(match) %>%
select(-estimand, -sample) %>%
kbl(booktabs = T, linesep = "") %>%
kable_styling(position = "center") %>%
kable_styling(latex_options = c("striped", "hold_position"))
```

estimator	est	se	t.stat	p
SATE	-5.448863	1.646936	-3.3084850	0.0009380
PATE	-5.448863	1.652232	-3.2978811	0.0009742
SATT	-1.277287	1.683284	-0.7588067	0.4479682
PATT	-1.277287	1.695820	-0.7531973	0.4513314
SATC	-7.016781	1.965677	-3.5696503	0.0003575
PATC	-7.016781	1.969424	-3.5628594	0.0003668

3 Appendix: Replicating Stata's Robust Standard Errors

Custom function by Jorge Cimentada that is used to replicate the robust standard errors in Stata:¹

```
robustse <- function(x, coef = c("logit", "odd.ratio", "probs")) {  
  suppressMessages(suppressWarnings(library(lmtest)))  
  suppressMessages(suppressWarnings(library(sandwich)))  
  
  sandwich1 <- function(object, ...) {  
    sandwich(object) *  
    nobs(object) / (nobs(object) - 1)  
  }  
  # Function calculates SE's  
  mod1 <- coeftest(x, vcov = sandwich1)  
  # apply the function over the variance-covariance matrix  
  
  if (coef == "logit") {  
    return(mod1) # return logit with robust SE's  
  } else if (coef == "odd.ratio") {  
    mod1[, 1] <- exp(mod1[, 1]) # return odd ratios with robust SE's  
    mod1[, 2] <- mod1[, 1] * mod1[, 2]  
    return(mod1)  
  } else {  
    mod1[, 1] <- (mod1[, 1] / 4) # return probabilities with robust SE's  
    mod1[, 2] <- mod1[, 2] / 4  
    return(mod1)  
  }  
}
```

¹<https://cimentada.jgithub.io/blog/2016-09-19-obtaining-robust-standard-errors-and-odds-ratios/obtaining-robust-standard-errors-and-odds-ratios-for-logistic-regression-in-r/>